

Lesson 11: Image processing

G Stefanescu — University of Bucharest

Parallel & Concurrent Programming
Fall, 2014

Generalities

Generalities:

- an *image* is represented as a two-dimensional array of *pixels* (picture elements)
- each pixel is either
 - a value in the *gray-scale* (usually, a number between 0 and 255)or it represents
 - a *color* (in rgb coding, i.e., the intensity of each primary red/green/blue color is specified; usually they are in the range 0-255)






Image processing is computationally intensive

Estimation (the need for parallel processing):

- suppose an image has 1024×1024 8-bit pixels (1Mb)
- if each pixel requires one operation, then 10^6 operations are needed for one frame;
- with computers performing 10^8 operations/sec this would take 10^{-2} s (= 10ms)
- in real-time applications usually the rate is 60-80 frames/sec.; hence each frame need to be processed in 12-16ms
- however, many image processing operations are complex, requiring much more than 1 operation per pixel; hence parallel processing may be useful here



Image processing

Image processing methods: these methods use and modify the pixels; most of them are very suitable for parallel processing; examples:

- basic low-level image processing, including noise cleaning or noise reduction, contrast stretching, smoothing, etc.
- *edge detection*
- *matching* an image against a given template
- *Hough transformation* (identify the pixels associated to straight lines or curves)
- *(fast) Fourier transform* - passing from an image to a frequency domain (and back)



Point processing

Point processing: methods acting on individual pixels (they are embarrassingly parallel)

- *Thresholding*: the pixels below a threshold are reduced to 0:

$$if(x < threshold) x = 0; \text{ else } x = 1$$

- *Contrast stretching*: the range of gray level values is extended to make details more visible

$$x = (x - x_l) \left(\frac{x_H - x_L}{x_h - x_l} \right) + x_L$$

a pixel of value x within the range $[x_l, x_h]$ is stretched to the range $[x_L, x_H]$ (this is often used in medical images, either for soft tissue portions, or for dense bone-like structures)



Histogram

Histogram

- an histogram shows the number of pixels in the image for each gray level (say, between 0 and 255)
- sequential code:

```
for (i=0; i<height_max; i++)  
    for (j=0; j<width_max; j++)  
        hist[p[i][j]] = hist[p[i][j]] + 1;
```

(hist[k] holds the number of pixels of gray level k)

- a parallel version may be easily developed (parallel addition of a set of numbers)



Smoothing, sharpening, noise reduction

Definitions:

- *smoothing* - suppress large fluctuations in intensity over the image area (can be achieved by reducing the high-frequency content)
- *sharpening* - accentuate the transitions, enhancing the detail (can be achieved in two ways: reduce low-frequency content or accentuate changes through differentiation)
- *noise reduction* - suppress a noise signal present in the image (not easy to distinguish noise from useful signal; a different method may be to capture the image more times and take the average on each pixel)



..Smoothing, sharpening, noise reduction

- the algorithms often require *local operations*, e.g., accessing all the pixels around the pixel to be updated

x_0	x_1	x_2
x_3	x_4	x_5
x_6	x_7	x_8

- Example: *mean*

$$x'_4 = \frac{x_0 + x_1 + x_2 + x_3 + x_4 + x_5 + x_6 + x_7 + x_8}{9}$$

it is used as a smoothing technique

- a sequential code for mean requires 9 operations for each pixel; hence sequential time complexity is $O(n)$



Parallel mean computation

Horizontal directions:

- 1 each processor (i, j) receives the value $x_{i,j-1}$ from *left* and adds it to an accumulating sum (the original value $x_{i,j}$ is retained)
- 2 then it receives the original value $x_{i,j+1}$ from *right* and adds it to the accumulating sum;

Hence processor (i, j) holds $x_{i,j} + x_{i,j-1} + x_{i,j+1}$; replace original values by these sums and repeat 1,2 for the vertical directions

- 3 processor (i, j) receives $x_{i-1,j} + x_{i-1,j-1} + x_{i-1,j+1}$ from above and adds it (retaining its previous sum $x_{i,j} + x_{i,j-1} + x_{i,j+1}$)
- 4 then it receives the sum $x_{i+1,j} + x_{i+1,j-1} + x_{i+1,j+1}$ from below and adds it getting the final sum $x_{i-1,j} + x_{i-1,j-1} + x_{i-1,j+1} + x_{i,j} + x_{i,j-1} + x_{i,j+1} + x_{i+1,j} + x_{i+1,j-1} + x_{i+1,j+1}$



Median

- mean value method tends to blur edges or other sharp details
- an alternative method is to use the *median*:
 - order the values of the neighborhood pixels and
 - choose the center value (provided an odd number of cells are compared)
- for the described 3×3 structure x_0, \dots, x_8 , order the values as $v_0 \leq v_1 \leq \dots \leq v_8$ and choose v_4
- one may use bubble sort, but stop when the center value was found (after 5 steps); this requires $8 + 7 + \dots + 4 = 30$ comparisons, hence $30n$ operations (for n pixels)



Parallel code / for median

- One may use a mash sorting algorithm, e.g., shear-sort;
- For greater speed an approximation method may be used:
 - use compare-and-exchange to sort any row:
 - Stage 1: $x_{i,j-1} \longleftrightarrow x_{i,j}$
 - Stage 2: $x_{i,j} \longleftrightarrow x_{i,j+1}$
 - Stage 3: $x_{i,j-1} \longleftrightarrow x_{i,j}$
 - repeat for columns:
 - Stage 1: $x_{i-1,j} \longleftrightarrow x_{i,j}$
 - Stage 2: $x_{i,j} \longleftrightarrow x_{i+1,j}$
 - Stage 3: $x_{i-1,j} \longleftrightarrow x_{i,j}$
 - the value $x_{i,j}$ after these 6 steps may not be the median, but it is a good approximation

This is the basic code for cell (i, j); it interferes with the code for other cells.



Weighted masks

- mean method gives equal weights to all neighborhood pixels
- generally, a *weighted mask* may be used
- for our standard 3×3 structure x_0, \dots, x_8 and weights w_0, \dots, w_8 , the new center pixel value is

$$x'_4 = \frac{w_0x_0 + w_1x_1 + w_2x_2 + w_3x_3 + w_4x_4 + w_5x_5 + w_6x_6 + w_7x_7 + w_8x_8}{k}$$

- k is used to maintain a correct gray-scale balance; usually k is $\sum_{i=0}^8 w_i$
- this operation may be seen as the “cross-correlation” of vectors x and w
- masks of other size may also be used; e.g., 5×5 , 9×9 , etc.



..Weighted masks

Examples (of masks):

- | | | |
|---|---|---|
| 1 | 1 | 1 |
| 1 | 1 | 1 |
| 1 | 1 | 1 |

 and $k = 9$ — may be used to compute mean

- | | | |
|---|---|---|
| 1 | 1 | 1 |
| 1 | 8 | 1 |
| 1 | 1 | 1 |

 and $k = 16$ — a noise reduction mask

- | | | |
|----|----|----|
| -1 | -1 | -1 |
| -1 | 8 | -1 |
| -1 | -1 | -1 |

 and $k = 9$ — a sharpening filter mask



Edge detection

Edge detection:

- object identification often requires to find *edges*
- an “edge” is a *significant change* of the gray level intensity

Suppose a one variable function $f(x)$ is considered (e.g., corresponding to a row);

- if f is differentiable, then:
 - its derivative $\partial f / \partial x$ has a *spike* when f has a significant change
 - the polarity of this spike gives the sense of the changing: positive (resp. negative) \Rightarrow increasing (resp. decreasing)
- if f is double differentiable, then its second derivative has a *zero* in the interval where f has such a significant change



..Edge detection

Suppose a (two-dimensional) image is considered; then the change in gray level have a

- *gradient magnitude* (number)

$$\nabla f = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$

- ... and a *gradient direction* (the angle with respect to the y-axis)

$$\phi(x, y) = \tan^{-1} \left(\left(\frac{\partial f}{\partial y}\right) / \left(\frac{\partial f}{\partial x}\right) \right)$$

These formulas may be used to identify the edges in an image (the 1st one is approximated as $\nabla f = \left|\frac{\partial f}{\partial x}\right| + \left|\frac{\partial f}{\partial y}\right|$ to simplify the computation)



Edge detection masks

- as usual, for discrete functions derivatives are approximated by

differences: for a 3×3 group,

x_0	x_1	x_2
x_3	x_4	x_5
x_6	x_7	x_8

we have:

$$\frac{\partial f}{\partial x} \approx x_5 - x_3 \quad \text{and} \quad \frac{\partial f}{\partial y} \approx x_7 - x_1$$

hence

$$\nabla f = |x_5 - x_3| + |x_7 - x_1|$$

- to compute this,
 - two masks may be used: one for $x_5 - x_3$, one for $x_7 - x_1$;
 - then add the resulting absolute values(the computation for each mask may be made in parallel)



Prewitt operator

Prewitt operator uses more points to approximate the gradient:

$$\frac{\partial f}{\partial x} \approx (x_2 - x_0) + (x_5 - x_3) + (x_8 - x_6)$$

$$\frac{\partial f}{\partial y} \approx (x_6 - x_0) + (x_7 - x_1) + (x_8 - x_2)$$

Then,

$$\nabla f = |x_2 - x_0 + x_5 - x_3 + x_8 - x_6| + |x_6 - x_0 + x_7 - x_1 + x_8 - x_2|$$

which, as above, requires two masks (one for each module)

-1	0	1
-1	0	1
-1	0	1

-1	-1	-1
0	0	0
1	1	1



Sobel operator

Sobel operator is a popular edge detection method; it uses a different approximation method for the gradient:

$$\frac{\partial f}{\partial x} \approx (x_2 - x_0) + 2(x_5 - x_3) + (x_8 - x_6)$$

$$\frac{\partial f}{\partial y} \approx (x_6 - x_0) + 2(x_7 - x_1) + (x_8 - x_2)$$

which, as above, requires two masks (one for each module)

-1	0	1
-2	0	2
-1	0	1

-1	-2	-1
0	0	0
1	2	1

Usually, the operators based of the 1st-order derivatives enhance noise; but Sobel operator has also a smoothing action.



Laplace operator

Laplace operator uses the 2nd-order derivatives

$$\nabla^2 f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$$

approximated to

$$\nabla^2 f = 4x_4 - (x_1 + x_3 + x_5 + x_7)$$

and computed using a single mask

0	-1	0
-1	4	-1
0	-1	0

Notice: We have studied this operator in other lectures, e.g., for heat distribution problem.



Edge detection

An original image



The effect of Sobel and Laplace operators

Sobel →



Laplace →






Hugh transform

Hugh transform:

- a method to “*find the parameters of equations of lines that most likely fit sets of pixels in an image*”
- may be used to fill in the gap between the points obtained in edge-detection result
- a line is described by the equation $y = ax + b$
- a direct search of the line including the largest number of pixels from a set of n pixels is expensive: $O(n^3)$ (for any 2 points check how many other points are on their line)
- rearrange the equation as $b = -xa + y$; then all points (x_i, y_i) on the line have the same associated (a, b) pair
- finally, count the number of points mapped to an (a, b) pair

..Hugh transform





Find “most likely” lines

1st version: Cartesian coordinates


- a single original point (x_1, y_1) is mapped into *all* (a, b) -points of the line

$$b = -x_1 \cdot a + y_1$$

- a discrete grid for pairs (a, b) is used; the computation is then rounded to the nearest possible (a, b) coordinate
- this is to be repeated for all given points (x_i, y_i)
- each pair (a, b) uses an accumulator to count the number of points mapped into it
- finally, the point (a, b) with the maximum value is chosen

Disadvantage: can not handle vertical lines.

..Hugh transform





..Find “most likely” lines

2nd version: Polar coordinates

- a single original point (x_1, y_1) is mapped into *all* (r, θ) -points satisfying

$$r = x_1 \cos\theta + y_1 \sin\theta$$

- the rest of the procedure is as in the previous case:
 - a grid of points (r, θ) is selected,
 - the number of pixels mapped into each (r, θ) -point is counted, and
 - the (r, θ) -point with the maximum value is selected
- it works well for all directions



Implementation

Implementation (using the standard image representation, i.e., origin = top/left corner):

- the parameter space is divided into small rectangular regions
- each region has an associated accumulator
- accumulators for the regions were a pixel maps into are incremented
- if k intervals are chosen for θ , then the complexity is $O(kn)$
- the complexity can be significantly reduced by limiting the range of lines for individual pixels using some criteria



..Implementation

Sequential code: (only one θ is used, based on the gradient function)

```
for (x=0; x<xmax; x++)  
    for (y=0; y<ymax; y++) {  
        sobel(&x,&y,dx,dx) /* find x,y gradients */  
        magnitude = grad_max(dx,dy);  
        if (magnitude > threshold) {  
            theta = grad_dir(dx,dy) /* use atan() fn */  
            theta = theta_quantize(theta);  
            r = x * cos(theta) + y * sin(theta);  
            r = r_quantize(r);  
            acc[r][theta]++;  
            append(r, theta, x, y); /* gather points */  
        }  
    }  
}
```



..Implementation

- from the resulting matrix `acc[][]` the points `(r, theta)` realizing a (local) maximum are chosen (hence an optimization algorithm has to be used/implemented)

Parallel code:

- there is a lot of room for parallelization in the above sequential algorithm, e.g.:
 - the accumulators may be computed in parallel
 - they use the same image, hence a shared memory model may be selected
 - only read actions on the image are performed, hence no critical sections are necessary



Transformation into frequency domain

Fourier transform:

- very useful transformation, used in many areas of science and engineering
- in image processing it was successfully applied for *image enhancement*, *restoration*, and *compression*
- we start with the one-dimensional case:
 - a periodic function $x(t)$ (of time) can be decomposed into a series of sinusoidal waveforms of various frequencies and amplitudes;
 - for each frequency f one gets an associated value $X(f)$
- the resulting series is called *Fourier series*; the transform is called *Fourier transform*



Fourier series

Fourier series:

- a *Fourier series* is a summation of sine and cosine terms

$$x(t) = \frac{a_0}{2} + \sum_{j=1}^{\infty} \left(a_j \cos \left(2\pi \frac{j}{T} t \right) + b_j \sin \left(2\pi \frac{j}{T} t \right) \right)$$

- T is the period of $x(t)$; $1/T = f$ is the frequency
- a more convenient representation (using complex numbers) is

$$x(t) = \sum_{j=-\infty}^{\infty} X_j e^{2\pi i \frac{j}{T} t}$$

X_j is called the j -th Fourier coefficient; $i = \sqrt{-1}$



Fourier transform

Fourier transform (for continuous functions):

- (direct) *Fourier transform*: given a continuous function of time $x(t)$, the function on frequency $X(f)$, called the *spectrum* of $x(t)$, is defined by

$$X(f) = \sum_{-\infty}^{\infty} x(t) e^{-2\pi i f t} dt$$

- *inverse Fourier transform*: given a continuous function of frequency $X(f)$, the function on time $x(t)$ is defined by

$$x(t) = \sum_{-\infty}^{\infty} X(f) e^{2\pi i f t} df$$

- key result: direct and inverse Fourier transforms are mutually converse one to the other



Discrete Fourier transform

Fourier transform (for discrete functions): similar, but integrals are replaced by finite sums

- *discrete Fourier transform* (DFT):

$$X_k = \frac{1}{N} \sum_{j=0}^{N-1} x_j e^{-2\pi i \left(\frac{jk}{N} \right)}$$

- *inverse Fourier transform*:

$$x_k = \frac{1}{N} \sum_{j=0}^{N-1} X_j e^{2\pi i \left(\frac{jk}{N} \right)}$$

- $0 \leq k < N$; N (real) numbers x_0, \dots, x_{N-1} produce N (complex) numbers X_0, \dots, X_{N-1}

The factor $1/N$ will be mostly omitted in the sequel; however, it has to be finally inserted to get a proper result



..Discrete Fourier transform

Example (16 points x_0, \dots, x_{15}):

$$X_0 = \frac{1}{16} (x_0 e^{-2\pi i (0 \frac{0}{16})} + x_1 e^{-2\pi i (1 \frac{0}{16})} + \dots + x_{15} e^{-2\pi i (15 \frac{0}{16})})$$

$$X_1 = \frac{1}{16} (x_0 e^{-2\pi i (0 \frac{1}{16})} + x_1 e^{-2\pi i (1 \frac{1}{16})} + \dots + x_{15} e^{-2\pi i (15 \frac{1}{16})})$$

$$X_2 = \frac{1}{16} (x_0 e^{-2\pi i (0 \frac{2}{16})} + x_1 e^{-2\pi i (1 \frac{2}{16})} + \dots + x_{15} e^{-2\pi i (15 \frac{2}{16})})$$

$$X_3 = \frac{1}{16} (x_0 e^{-2\pi i (0 \frac{3}{16})} + x_1 e^{-2\pi i (1 \frac{3}{16})} + \dots + x_{15} e^{-2\pi i (15 \frac{3}{16})})$$

\vdots

$$X_{15} = \frac{1}{16} (x_0 e^{-2\pi i (0 \frac{15}{16})} + x_1 e^{-2\pi i (1 \frac{15}{16})} + \dots + x_{15} e^{-2\pi i (15 \frac{15}{16})})$$



..Discrete Fourier transform

In a marticial form, this transformation may be written as follows
(using $w = e^{-2\pi i(\frac{1}{16})}$):

$$\begin{pmatrix} X_0 \\ X_1 \\ X_2 \\ X_3 \\ \vdots \\ X_{15} \end{pmatrix} = \frac{1}{16} \begin{pmatrix} 1 & 1 & 1 & 1 & \dots & 1 \\ 1 & w & w^2 & w^3 & \dots & w^{15} \\ 1 & w^2 & w^4 & w^6 & \dots & w^{2 \cdot 15} \\ 1 & w^3 & w^6 & w^9 & \dots & w^{3 \cdot 15} \\ \vdots & & & & & \\ 1 & w^{15} & w^{15 \cdot 2} & w^{15 \cdot 3} & \dots & w^{15 \cdot 15} \end{pmatrix} \cdot \begin{pmatrix} x_0 \\ x_1 \\ x_2 \\ x_3 \\ \vdots \\ x_{15} \end{pmatrix}$$



Fourier transform in image processing

Two-dimensional Fourier transform (the factor $1/(NM)$ is omitted)

- a 2-dim DFT is

$$X_{lm} = \sum_{j=0}^{N-1} \sum_{k=0}^{M-1} x_{jk} e^{-2\pi i \left(\frac{jl}{N} + \frac{km}{M} \right)}$$

where j (resp. k) is the row (resp. column) coordinate

- the formula may be rewritten as

$$X_{lm} = \sum_{j=0}^{N-1} \left[\sum_{k=0}^{M-1} x_{jk} e^{-2\pi i \left(\frac{km}{M} \right)} \right] e^{-2\pi i \left(\frac{jl}{N} \right)}$$

showing that a 2-dim DFT may be obtained in two phases:

- an (inner) 1-dim DFT operating on row, followed by
- a 1-dim DFT operating on columns



..Fourier transform in image processing

Application of DFT in image processing:

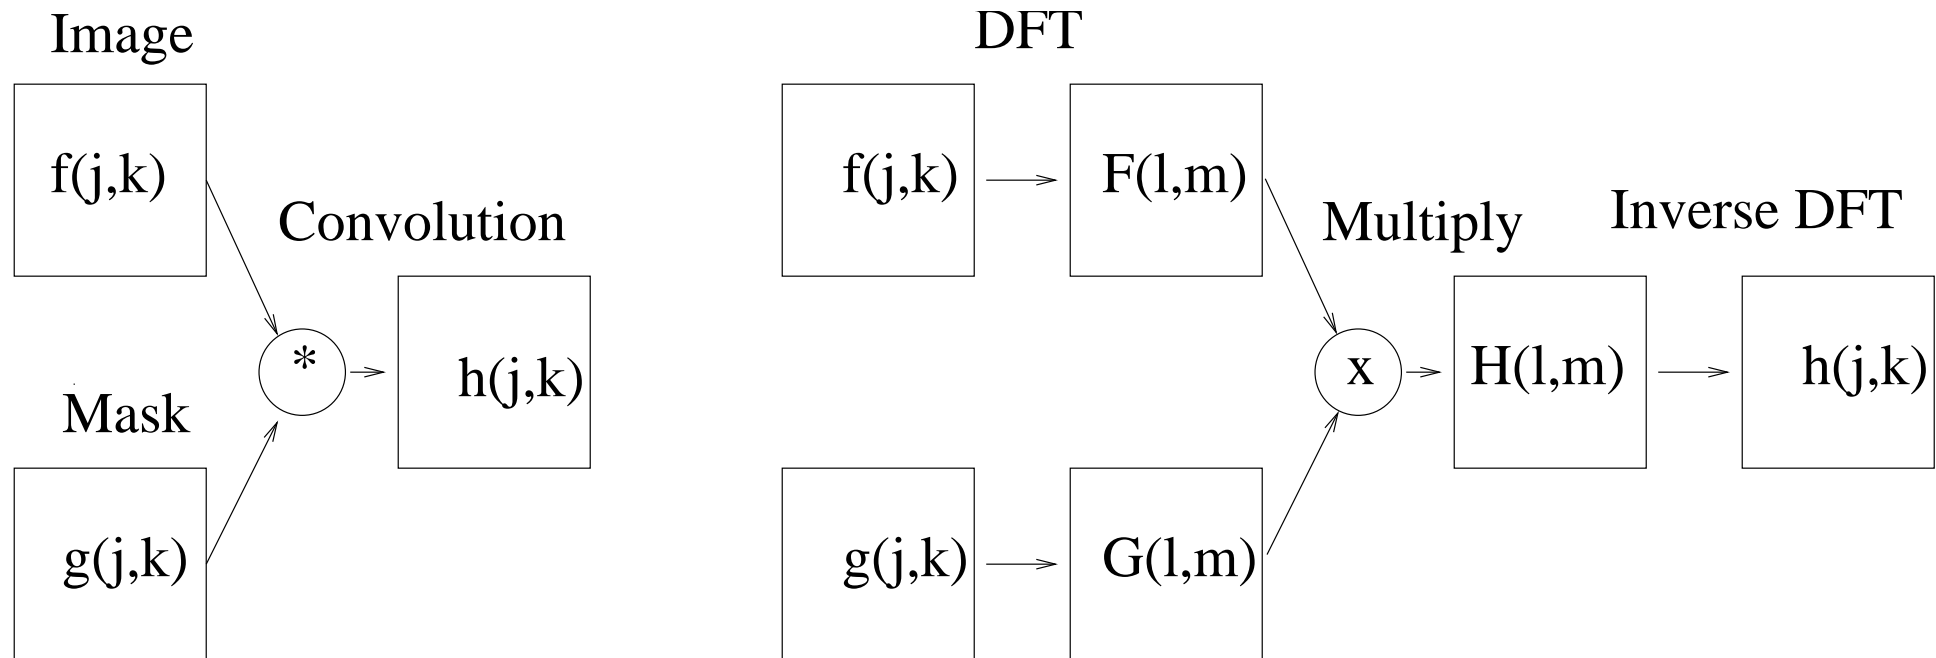
Frequency filtering is used for both smoothing and edge detection
—earlier we have used weighted masks for that purpose;
—but this is just a particular case of *convolution*

$$h(i, j) = g(j, k) * f(j, k)$$

where $g(j, k)$ describes the mask and $f(j, k)$ the image
—the convolution of functions corresponds to the product of their
Fourier transform (“ \times ” denotes element-wise multiplication)

$$H(l, m) = G(l, m) \times F(l, m)$$

..Fourier transform in image processing





Sequential code for DFT

Sequential code: denote $w = e^{-2\pi i/N}$; then $X_k = \sum_{j=0}^{N-1} x_j (w^k)^j$

```
for (k=0; k<n; k++) {  
    X[k] = 0;  
    a = 1;  
    for (j=0; j<N; j++) {  
        X[k] = X[k] + a * x[j];  
        a = a * wk;  
    }  
}
```



Parallel DFT

Parallel implementation (direct approach):

- use a master/slave approach (one slave for computing each $X[k]$); with N processes this gives an $O(N)$ algorithm

Pipeline implementation: Unfolding the inner loop we get

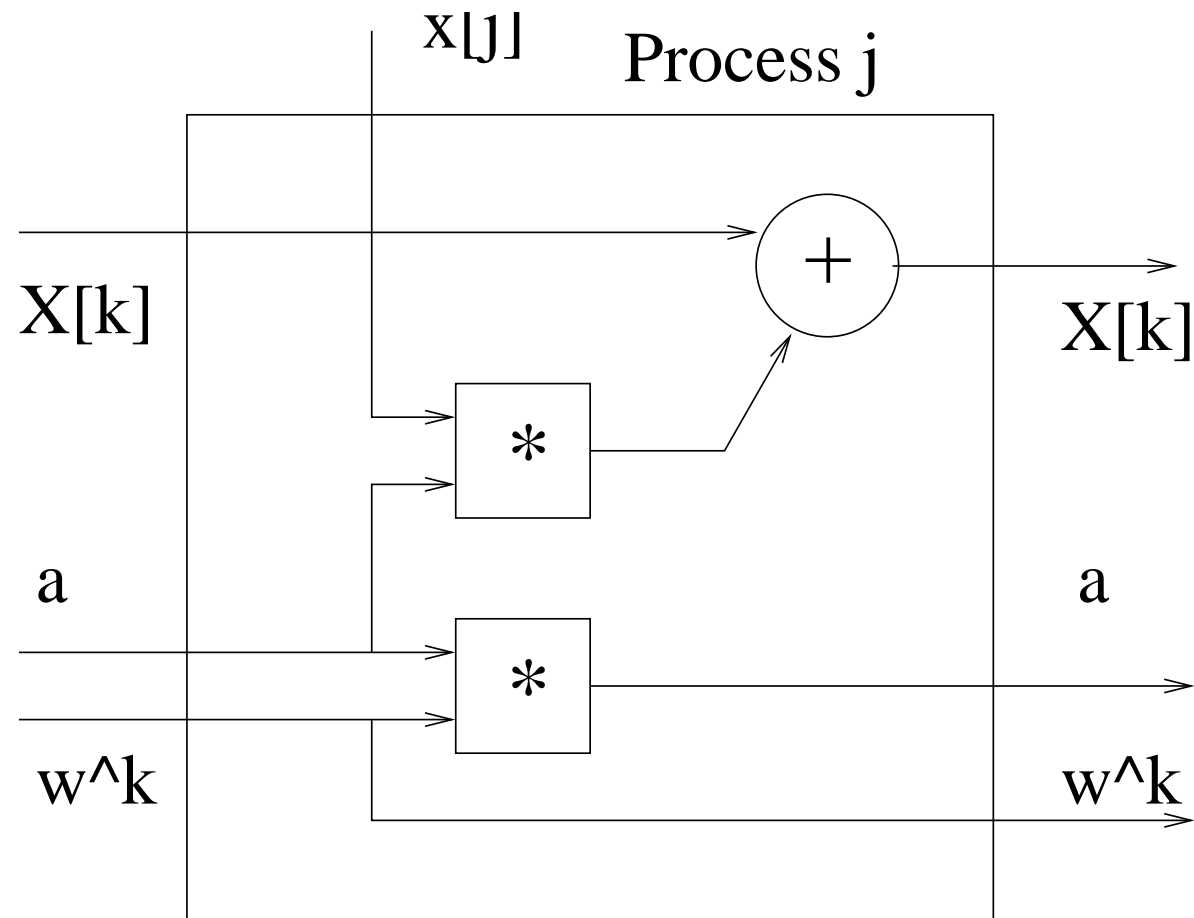
```
X[k] = 0;  
a = 1;  
X[k] = X[k] + a * x[0];  
a = a * wk  
X[k] = X[k] + a * x[1];  
a = a * wk  
⋮
```



..Parallel DFT

Pipeline implementation:

- one stage of the pipeline implementation is





Fast Fourier transform (FFT)

Fast Fourier transform: it's a method to reduce (sequential) time complexity from $O(N^2)$ to $O(N \log(N))$

Based on a recursive procedure:

- start with

$$X_k = \frac{1}{N} \sum_{j=0}^{N-1} x_j w^{jk}$$

- divide the summation into two parts

$$X_k = \frac{1}{N} \left[\sum_{j=0}^{(N/2)-1} x_{2j} w^{2jk} + \sum_{j=0}^{(N/2)-1} x_{2j+1} w^{(2j+1)k} \right]$$



..FFT

- rewrite the sum as $X_k = (1/2)[X_{even} + w^k X_{odd}]$, namely

$$X_k = \frac{1}{2} \left[\frac{1}{N/2} \sum_{j=0}^{(N/2)-1} x_{2j} w^{2jk} + w^k \frac{1}{N/2} \sum_{j=0}^{(N/2)-1} x_{2j+1} w^{2jk} \right]$$

- notice that $w^{k+(N/2)} = e^{(-2\pi i/N)(k+(N/2))} = -e^{(-2\pi i/N)k} = -w^k$
hence


$$X_{k+(N/2)} = (1/2)[X_{even} - w^k X_{odd}]$$

showing that X_k and $X_{k+(N/2)}$ could be computed using two transforms involving $N/2$ points

- the procedure can now be recursively applied
- a sequential code leads to an $O(N \log(N))$ algorithm

..FFT

The recursive structure of *parallel FFT* is:



Comments:

- notice that the terms w^k used in the figure are dependent on the number of points N , say w_N^k
- hence in the inner boxes of $N/2$ points different values have to be used $w_{N/2}^k$ for $k = 0, \dots, N/2 - 1$
- ... but we can normalize them: $w_{N/2}^k = w_N^{2k}$



..Parallel FFT

Analysis:

- Computation: with p processes and N points, each process will compute N/p points; each points requires 2 operations; with $\log(N)$ steps this gives $t_{comp} = 2(N/p) \log(N) = O(N \log(N))$
- Communication: if $p = N$, communication occurs at each step and one data is exchanges between pair of processors (a finer analysis may be give, depending on the network structure of the processors)